


PAVS Model Lifecycle & Delivery Strategy 2026

4-Phase Model Lifecycle



Team Core Capabilities:

- **Model Enablement**: Off-the-shelf installation, dependency resolution, initial benchmarking
- **Performance Profiling**: Compute analysis, hotspot identification, bottleneck detection
- **Benchmarking & Validation**: Task metrics, Identify/suggest right model per domain/client, competitive analysis (GPU vs CPU vs NPU), CI/CD automation
- **Model Optimization**: Quantization (INT8, FP16), kernel optimization
- **Training Infrastructure**: Custom model training, domain-specific adaptation, fine-tuning on custom datasets
- **Deployment & Release**: Docker containerization, HuggingFace distribution, automated documentation

2026 Release Timeline

R1: Foundation

- Complete lifecycle framework for all models
- Streamline release logistics & CI/CD
- Deploy automated docs ([GitHub Pages/AMD docs](#))

R2: Scale

- Single SDK installer (Nvidia/Intel style)
- 8 production-ready models (P0/P1)
- Automated benchmarking pipeline
- Quantization & Infra Beta (vLLM, ONNX-RT, Llama.cpp, Lemonade)
- Sample apps (ChatBot, Object Detection, GStreamer+ROCm)

R3: Expand

- 8 more release-ready models
- Enhanced quantization & infrastructure
- Model deployment & robotics simulator
- Scale fine-tuning operations

R4: Complete 1st cycle

- 2 additional production models (20 total)
- Multi-segment deployment (Robotics, Healthcare, Industrial)
- Model training infrastructure
- 2027 strategy planning

Open Source Release & Distribution Strategy

HuggingFace Hub Distribution

What We Publish:

- 🤖 **Pre-optimized Model Weights**: ONNX, TensorRT, ROCm-ready
- 📦 **Quantized Variants**: INT8, FP16 for edge deployment
- 🐳 **Docker Images**: GPU/NPU/Hybrid ready-to-run containers
- 📊 **Benchmark Results**: Verified performance metrics included
- 📄 **Model Cards**: Documentation, usage examples, hardware requirements

HuggingFace Advantages:

- ✅ **Public Discovery**: Searchable by AI/robotics community
- ✅ **Version Control**: Track model iterations & improvements
- ✅ **Direct Download**: `pip install` or one-click deployment
- ✅ **Community Engagement**: Issues, feedback, contributions
- ✅ **Ecosystem Integration**: Works with Transformers, ONNX-RT, vLLM

Infrastructure Tools (Open Source Ready):

- **vLLM**: Efficient LLM serving & inference
- **Llama.cpp**: Optimized CPU/GPU inference
- **ONNX Runtime**: Cross-platform model deployment
- **Lemonade Profiler**: Performance analysis tool