

Why a Model Zoo?

From a 100+ GB Tarball to an Off-the-Shelf Product Experience

What PAVS needs from a Model Zoo — and what we must not turn it into

1 · The Off-the-Shelf Promise

What an application developer actually wants

- Download a model → **start building the application today**
- No knowledge of layers, kernels, quantization, or backends
- Just a **prepackaged model, best-suited to the task**, ready to run on AMD hardware

The developer knows the **inputs and outputs** of the model and builds his product around them — nothing more.

The experience we are selling

Pick a task — detect / classify / segment / LLM



Download the packaged model — already optimized



Build the product — runs on CPU · GPU · NPU

Without a zoo, the developer gets weights and a to-do list. **With a zoo, he gets a model he can ship today.**

2 · Two Kinds of Users — Two Kinds of Products

AIG Git + Model Zoo serves...

The Research Engineer

- Knows the **nitty-gritty** of the model internals
- Rewires layers, swaps operators, debugs graphs
- Cares *which layer breaks, which workflow is GPU-heavy, which lighter variant to substitute*
- Works **inside** the model

Audience: people who build and modify models.

PAVS serves...


The Application Engineer

- Uses AI models as **black boxes**
- Knows only the **inputs and outputs**
- Does **not** need to know which layer breaks or which variant is lighter
- Builds the **application around** the model

Audience: people who ship products using models.

Our zoo brings a **different kind of user** onto AMD silicon — and needs to be shaped for **them**, not for the researcher.


3 · NVIDIA's Blueprint: NGC ↔ Isaac ROS



NGC Catalog — curated, prepackaged models (PeopleNet, Mistral-7B Int4, ...)

The pattern that makes it work

- **NGC Catalog** hosts the ready-to-run models
- **Isaac ROS** consumes them directly into robotics pipelines
- The developer pulls a model and it *just runs* — no export, no tuning



AMD needs exactly this pairing: a **zoo** that hosts optimized models + an **SDK / robotics stack** that consumes them off-the-shelf.

4 · The Tarball Problem

How we ship today

- **Every model** is packaged into the release tarball
- Add optimized variants of CNNs **and** LLMs and it balloons fast
- Realistically **100+ GB** — and growing with every new variant

Every user carries **the entire zoo on their disk** — even if they only ever use a couple of CNNs.

Monolithic Tarball · 100+ GB

YOLOv12 YOLO26 fp16/int8 MobileSAM CenterPoint Llama-3 int4
SmolVLA + every variant


↓ forced onto every disk

User who only wanted **2 CNNs** still pays the full 100+ GB


A zoo fixes this: **pull only the model you need**, keep the SDK lean.

5 · With the Zoo vs Without — Two Flows to the Same Model

AIG hands us a PyTorch model — the zoo stores every optimized artifact so the customer never has to




✓ With the Zoo — one download



Every optimized artifact is **ready to pull** — the heavy lifting was done once, by us.

✗ Without the Zoo — run every stage




Every customer re-runs the **whole toolchain on their own disk** — hours of work, before line one of the app.

6 · The Current PAVS Model Zoo — Live Today

Built close to the Qualcomm model — the structure is live and already serving models

Live now on Hugging Face

- **AMD-PAVS-AI** — 49 models, per-task, ready to pull
- Structured **close to Qualcomm AI Hub** — per-backend artifacts
- Graphics / dashboard not there yet — but the **storage + structure is live and serving**



AMD-PAVS-AI on Hugging Face — 49 models, per-task, ready to pull

This is the pattern **NVIDIA (NGC)** and **Qualcomm (AI Hub)** already ship — pre-baked, ready-to-run models per target. PAVS is on the same path.

7 · What PAVS Expects From a Model Zoo

Just a storage space — like Hugging Face

- We **don't care who owns or maintains it** — could be AMD central-wide infrastructure
- One requirement: it **hosts weights and serves them to customers**

What we do need: operational flexibility

- Freedom to **create directories** / structure as our workflows demand
- **Store the weights** for every optimization stage
- **Make them public** for customer consumption on our schedule

Model Zoo = Storage + Freedom


-  Create our own directory structure
-  Store weights for every stage
-  Publish for customers directly
-  Control what is public, and when

Think **Hugging Face for AMD**: a backing store we operate freely — not a system we must conform to.

The Founding Promise — Research Desk → Vertical Software

1. PAVS transforms raw researcher code into production-ready vertical software

What we receive: Research Repo



What we receive: Research Code

```
# rocm-scripts/test/pytorch/yolo12n.py
# .....
# .....
# .....
# .....
# .....
# .....
recalled = gt_labels & detected_classes
recall = len(recalled) / len(gt_labels)
sanity_pass = recall >= RECALL_THRESHOLD
assert sanity_pass, (
    f"recall {recall:.0%} < {RECALL_THRESHOLD:.0%}")

return lats, {
    "architecture": "CNN",
    "task": "object_detection",
    "num_detections": num_detections,
    "recall": round(recall, 4),
}
if __name__ == "__main__":
    sys.exit(run(run_benchmark))
```

Single script, no structure, no device flows

What we deliver: Vertical Software

```
examples/yolov12/
├── Makefile
├── METRICS_TABLE.md
├── README.md
├── app/
│   ├── routes/
│   ├── services/
│   └── main.py & config.py
├── scripts/
│   ├── benchmark.py
│   ├── export_onnx.py
│   ├── profile_model.py
│   ├── run_val.py
│   └── update_metrics_table.py
├── tests/
│   ├── test_api.py
│   ├── test_docker.py
│   ├── test_migraphx.py
│   └── test_yolo_workflow.py
├── weights/
├── datasets/
├── requirements.txt
├── requirements_cpu.txt
└── requirements_npu.txt
```

App, scripts, tests, per-device requirements — production-ready

Flat list of scripts — no structure, no device flows, no Make targets

Physical AI & Vertical Software: Researcher's Desk → Vertical Software — same commands, every device, production-ready.

8 · What We Should NOT Do

The initial promise: convert research code into **vertical software** — a zoo must not reverse that

The anti-pattern

If AIG **hosts** the zoo and mandates we place our models in it:

1. To appear on **their dashboard**, we must write research code that **fits their git repo**
2. We end up rebuilding a **research setup** — **not vertical software**
3. This turns our SDK into a mirror of **their repo and dashboard**

This direction **nullifies the value-add** that Physical AI & Vertical Software brings — the exact opposite of our founding promise.

The right shape

- Our job (per the SW-stack promise) is **Research Desk** → **Vertical Software**
- Fitting AIG's repo/dashboard is **book-keeping**, not a **customer-facing** deliverable

Vertical software — customer-facing, product-grade

▲ keep on top

AIG repo / dashboard sync — do it as *book-keeping*, if at all

If it must be done, do it as book-keeping — never let it turn the SDK into a customer-facing clone of the research repo.

The Ask

A model zoo that is **storage + operational freedom** — so
PAVS ships an **off-the-shelf, product-grade experience** on
AMD silicon

**Host the optimized artifacts · keep the SDK lean · serve the application engineer ·
stay vertical software**